

SENTINEL: Scalable Environmental Network for Temporal Intelligence and Ecological Learning

Multimodal Artificial Intelligence for Early Water Pollution Detection

Entry to the Stockholm Junior Water Prize (2026)

Bryan Cheng & Austin Jin

New York, USA

I. Abstract

Today's water-pollution monitoring is too slow and too sparse: a handful of aging sensors, spread thin across vast river networks, flagging contamination only after it spreads. The cost of these blind spots is measured in human lives and billions of dollars — from lead contamination to toxic algal blooms that shut off drinking water for whole cities. SENTINEL is the first multimodal deep-learning early-warning system for water contamination, fusing five independent streams of environmental data: water chemistry, satellite imagery, microbial DNA, gene-expression stress signals in aquatic organisms, and the movement patterns of aquatic animals. Each stream is read by a dedicated, custom-designed neural network that matches or exceeds the state-of-the-art, and when one stream flags an anomaly the others confirm and classify it. Trained entirely on public data, SENTINEL detected the Gulf of Mexico dead zone 87 days before NOAA's field survey and warned an average of 32 days ahead across 31 real contamination events. To screen widely at low cost, we designed and built SENTINEL-Lite: a custom carbon-fiber drone that surveys water bodies from the air for roughly \$0.50 per site using HydroDenseNet. The two systems work in tandem — Lite screens broadly, then hands the highest-risk sites to SENTINEL for deep multimodal confirmation.

II. Table of Contents

I. Abstract	1
II–VI. Preliminary Matters	2
1. Introduction	2
2. Materials and Methods	4
3. Results	8
4. Discussion	17
5. Conclusions	19
6. References	20

III. Key Words

Water quality monitoring; artificial intelligence; multimodal fusion; early warning system; anomaly detection; deep learning; environmental sensing; satellite remote sensing; metagenomics; behavioral ecotoxicology; conformal prediction; causal discovery; uncertainty quantification; risk assessment

IV. Abbreviations and Acronyms

AI (Artificial Intelligence), AUROC (Area Under the Receiver Operating Characteristic Curve), CI (Confidence Interval), CKA (Centered Kernel Alignment), CLR (Centered Log-Ratio), DO (Dissolved Oxygen), ECE (Expected Calibration Error), eDNA (Environmental DNA), EMP (Earth Microbiome Project), EPA (U.S. Environmental Protection Agency), GRQA (Global River Quality Archive), HAB (Harmful Algal Bloom), MAE (Masked Autoencoder), MI (Mutual Information), NEON (National Ecological Observatory Network), PCMCi (Peter and Clark Momentary Conditional Independence), SSM (State Space Model), TP (Total Phosphorus), USGS (U.S. Geological Survey), ViT (Vision Transformer), WQ (Water Quality).

V. Acknowledgements

We thank our research mentor for guidance on methodology and experimental design. We are grateful to the providers of public datasets: USGS (NWIS), EPA (WQP, ECOTOX), NEON, ESA (Sentinel-2), the Earth Microbiome Project, and the GRQA team. All computation was performed on a personal NVIDIA RTX 4060 GPU. We conducted all programming, model development, data analysis, and paper writing ourselves.

VI. Biography

Austin Jin and Bryan Cheng are juniors at William A. Shine Great Neck South High School in Great Neck, New York. They share interests in biology, environmental science, and public health, with a particular focus on how computational tools and artificial intelligence can address pressing challenges in ecological monitoring and human health. Their research together explores how multimodal machine learning can be applied to real-world water quality surveillance, with the broader goal of making environmental and public health monitoring more proactive, equitable, and accessible. Outside of research, both are active in their school and local community, and they plan to pursue studies at the intersection of biology, computation, and medicine in college.

1. Introduction

In 2014, the City of Flint, Michigan switched its water source without adequate treatment. Over 18 months, 100,000 residents were exposed to lead—yet agencies did not acknowledge the contamination until October 2015 [1]. The Flint crisis is not isolated: the 2015 Gold King Mine spill released 3 million gallons of toxic wastewater undetected, and the 2023 East Palestine derailment contaminated waterways

while monitoring systems remained silent [2]. Even EPA-documented events are frequently detected only after damage has occurred.

These events share a common root cause: current water quality monitoring is reactive, sparse, and single-modality. The US operates ~1,130 continuous monitoring stations across 3.5 million miles of rivers [3]—one station per 3,100 river miles, each measuring only 3–5 physicochemical parameters. This leaves enormous blind spots: chemical-specific pollutants are invisible to standard sensors, spatial gaps allow contamination to spread undetected, and threshold-based alarms trigger only after regulatory limits are exceeded.

Biology already runs a better detection system. Organisms respond at multiple timescales: fish alter behavior within minutes [4], dissolved oxygen shifts within hours, microbial communities reorganize over days [5], and satellite-observable properties change over weeks [6]. No existing system integrates these diverse signals. Prior AI focuses on single modalities—satellite WQ prediction [6], sensor foundation models [3], statistical biomonitors [4]—each with blind spots. None provide uncertainty quantification or coverage guarantees for regulatory adoption.

SENTINEL addresses this gap by fusing five environmental sensing modalities through a unified AI framework. The key insight is that different modalities respond to different contaminants at different timescales—fish change behavior within minutes of a toxic spill, microbial communities shift over days, and satellite imagery reveals algal blooms over weeks—and joint reasoning across all of them detects contamination earlier than any single modality alone.

This project makes nine contributions:

1. **SENTINEL-DB:** The largest multimodal water quality dataset ever assembled—over 390 million records from 19 public data sources spanning 105 countries.
2. **Five novel encoder architectures:** Each designed for its modality’s unique data characteristics, outperforming published SOTA where benchmarks exist (AquaSSM vs. MCN-LSTM, HydroViT vs. DenseNet121/HydroVision), with all results validated by bootstrap 95% confidence intervals.
3. **SENTINEL-Fusion:** A Perceiver IO-derived cross-modal temporal attention framework that achieves AUROC = 0.939 [0.922–0.956] and detects contamination events with calibrated uncertainty (ECE = 0.22).
4. **Real-data event validation:** 6 major contamination events detected on real downloaded USGS sensor data with mean lead time of 66 days, plus 6 real NEON events and 19 research-validated events spanning HABs, hypoxia, salinity intrusion, and agricultural runoff.
5. **Mathematical safety guarantees:** Conformal prediction provides distribution-free coverage guarantees (≥ 0.95) validated on 13,202 real encoder embeddings.
6. **Actionable deployment tools:** A composite risk index ranking 32 NEON sites, causal pollution pathway discovery, and cost-optimal sensor placement starting at \$0.50/site/year.
7. **Mechanistic and ecological interpretability:** Biologically-constrained encoders and causal discovery explain contamination rather than merely flagging it—recovering 375 causal pollution pathways (44 novel), tracing contaminant-specific molecular signatures, attributing microbial source, and linking water quality to keystone-species health and waterborne-disease risk.
8. **Ecosystem simulation:** A neural-ODE/physics digital twin forecasts ten coupled water-quality variables across six horizons (1–365 days, 45.5% lower error than physics-only) and supports counterfactual “what-if” intervention and restoration planning.
9. **Tiered deployment:** A drone/satellite SENTINEL-Lite screening layer surveys water bodies at ~\$0.50/site/year, and a learned cascade controller activates the expensive multimodal pipeline only where screening indicates risk—making nationwide coverage economically feasible.

The urgency of this work reflects a global crisis. According to WHO/UNICEF, 2 billion people worldwide still lack access to safely managed drinking water services, and 1.2 billion lack even basic service [18]. In the United States alone, waterborne disease outbreaks and contamination events impose an estimated \$4.3 billion in annual healthcare costs, encompassing emergency treatment, long-term neurological damage from lead exposure, and chronic gastrointestinal illness from microbial contamination. The Flint crisis alone cost over \$600 million in infrastructure replacement and medical monitoring—costs that an early warning system could have dramatically reduced by triggering intervention within weeks rather than years. Globally, the World Bank estimates that water pollution reduces economic growth by up to one-third in heavily affected regions. Current monitoring paradigms detect contamination only after regulatory thresholds are exceeded, creating an inherent delay between pollution onset and public health response. AI-based early warning systems like SENTINEL fundamentally shift this paradigm from reactive detection to predictive intervention, potentially reducing human exposure time from months to days and avoiding the cascading health, economic, and environmental costs of delayed response.

2. Materials and Methods

2.1 SENTINEL-DB: Dataset Construction

SENTINEL-DB consolidates 19 publicly available environmental data sources into the largest multimodal water quality dataset assembled (Table 1). Nine sources appear in Table 1; SENTINEL-DB additionally incorporates USGS BioData (701K biological records), Water Quality Portal cyanotoxin/nutrient (755K) and Canadian (787K) subsets, NOAA ERDDAP chlorophyll-a (146K), NHDPlusV2 stream topology (561 sites, 338 edges), and GBIF freshwater species occurrences (2,355).

Source	Modality	Records	Coverage
NEON Aquatic	Sensor	351.7M	34 US sites, 24 months
GRQA v1.3	Sensor	18.0M	94K sites, 105 countries
EPA WQP	Sensor	18.3M	18 HUC2 basins
USGS NWIS	Sensor	291K	1,130 stations
Sentinel-2 L2A	Satellite	2,986 tiles	Global water pixels
EMP 16S rRNA	Microbial	20,288	127 habitat types
NCBI GEO	Molecular	84K genes	4 transcriptomic datasets
EPA ECOTOX	Molecular	268K	1,391 chemicals
EPA ECOTOX Daphnia	Behavioral	29,421	Real conc.-response tests

Table 1. SENTINEL-DB data sources. Total: 390M+ records, ~85 GB.

All data sources are publicly accessible. NEON contributes the majority (351.7M rows of continuous sonde data at 15-minute resolution); GRQA provides geographic diversity (18M records, 94K sites, 105 countries) [7]; Sentinel-2 L2A provides spatially continuous observations every 5 days. Harmonizing required resolving differences in temporal resolution (15 min to monthly), spatial scale (point stations to 10 m pixels), units, and quality flags. The pipeline applies source-specific QC, standardizes to SI/UTC, and constructs co-registered training pairs. Constructing SENTINEL-DB required overcoming four major data harmonization challenges. First, temporal resolution mismatch: NEON sondes report every 15 minutes (96 readings per day), USGS NWIS provides instantaneous values at 15-minute to hourly intervals, EPA WQP contains grab samples at weekly to monthly frequency, and Sentinel-2 revisits every 5 days with cloud-dependent gaps. Aligning these required interpolation for sensor data (forward-fill with exponential decay weighting) and temporal binning for satellite-sensor co-registration (matching the closest cloud-free overpass within a 48-hour window to in-situ measurements). Second, spatial scale differences: point-based sensor stations measure water quality at a single location, while Sentinel-2 pixels represent 10-meter ground cells and EMP 16S samples integrate microbial communities over variable collection volumes. Co-registration used geographic buffering (500 m radius

for sensor-satellite pairs, watershed-level matching for microbial samples). Third, missing data handling: across 390 million records, approximately 12% contained at least one missing parameter. Rather than imputing missing values—which could introduce artificial patterns—the pipeline flags gaps and the fusion module handles missing modalities through confidence-weighted gating with renormalization at inference time. Fourth, quality control: each source required specific QC steps—NEON data was filtered by the finalQF flag (retaining only QF-passed records), USGS records were screened for instrument drift using rolling standard deviation thresholds, and EMP samples were rarefied to 10,000 reads to normalize sequencing depth. This four-stage pipeline—temporal alignment, spatial co-registration, missingness flagging, and source-specific QC—ensures that the training data reflects genuine environmental variation rather than measurement artifacts.





Fig 1. The five SENTINEL-DB training modalities and their fusion: (A) sensor water-quality time series, (B) satellite imagery with spectral profiles, (C) microbial OTU compositions, (D) molecular gene-expression heatmap, (E) *Daphnia* behavioral trajectories, and (F) Perceiver IO fusion of the five modality embeddings.

2.2 Modality-Specific Encoders

Each sensing modality requires a specialized neural architecture. SENTINEL’s pipeline operates in three stages: (1) five modality-specific encoders independently project their respective data types into a shared 256-dimensional embedding space; (2) a Perceiver IO fusion module with 64 learned latents performs cross-modal temporal attention across all available embeddings; and (3) four output heads produce calibrated anomaly probability, 8-class type classification, 8-class source attribution, and 5-tier cascade escalation. The total system comprises 65.3M parameters (62.3M encoders + 3.0M fusion), all trained on a single NVIDIA RTX 4060 (8 GB) in ~72 hours. The individual encoder architectures, each designed to address the unique data challenges of its modality, are described below.


AquaSSM (Sensor). Environmental sensor data presents two challenges: irregular sampling intervals (maintenance gaps, telemetry failures) and multi-timescale dynamics (acute spills unfold over hours, eutrophication over months). AquaSSM addresses both with a continuous-time SSM [8] that parameterizes state transitions as $h(t) = \exp(A \cdot \Delta t) h(t-1)$, naturally handling arbitrary time gaps. Eight



parallel channels span 1 hour to 365 days: short channels capture rapid conductance spikes, long channels track seasonal oxygen decline. Trained with masked-parameter pretraining on 127,421 real USGS NWIS sequences from 303 stations under a strict spatial holdout (62 test sites unseen during training; held-out reconstruction RMSE = 0.83).

Figure 2. AquaSSM sensor encoder. Eight parallel CT-SSM channels (1h–365d) process irregularly-sampled multi-parameter sensor data. Trained on 127,421 USGS sequences (303 stations, spatial holdout) (AUROC = 0.939).

HydroViT (Satellite). Satellite remote sensing fills the spatial gaps between point sensors, but paired satellite–in-situ training data is scarce and many pollutants are optically invisible. HydroViT v9 is a CNN-ViT hybrid: a 3-layer stride-1 CNN (10 → 32 → 64 → 128 channels) extracts local spectral features before a ViT-S/16 encoder captures global spatial relationships. Per-parameter band attention




weights are physics-initialized (e.g., chlorophyll-a prioritizes bands 2–4). Training proceeds in two phases: CNN+head training with frozen ViT (80 epochs), then full

Figure 3. HydroViT v9 satellite encoder. CNN-ViT hybrid with MAE pretraining on 2,986 Sentinel-2 patches and fine-tuning on 5,464 paired samples (water temp $R^2 = 0.893$).

fine-tuning (120 epochs) with cosine annealing. Trained on 5,464 co-registered Sentinel-2/in-situ pairs for 11 WQ parameters, HydroViT achieves water temperature $R^2 = 0.893$, outperforming DenseNet121 (HydroVision-style [23], $R^2 = 0.884$, +0.009) and excelling on TSS (+0.100) and phycocyanin (+0.097).

MicroBiomeNet (Microbial). Microbiome data is compositional—OTU abundances sum to a constant, creating spurious correlations if analyzed with standard methods



[11]. MicroBiomeNet operates natively in Aitchison simplex geometry: a CLR transform maps compositional data to Euclidean space, followed by an MLP and Neural ODE modeling temporal community dynamics. Trained on 20,288 EMP 16S rRNA samples spanning 127 habitats, it classifies aquatic sources into 8 categories (macro-F1 = 0.899 on a held-out EMP split, 0.831 under strict spatial holdout; per-class range 0.72–0.95).

ToxiGene (Molecular). Gene expression data is high-dimensional (~84K genes) but structured: genes




Figure 5. ToxiGene molecular encoder (145.2M parameters). Left: Hierarchical sparse architecture mirrors the biological gene → pathway → process → adverse outcome hierarchy, with sparse connections constrained by Reactome and AOP-Wiki. Right: Biologically-informed sparsity with 145.2M parameters constrained to mechanistically plausible relationships.

for multi-label toxicity classification directly from RNA-seq expression profiles. Closest prior art addresses different problems: MTForestNet [26] predicts chemical–gene interactions from molecular fingerprints (not transcriptomics), Senn et al. [27] cluster transcriptomic responses without contaminant labels, and Gagné et al. [28] use unsupervised pathway enrichment for effluent characterization.

BioMotion (Behavioral). Organismal behavior is the fastest biological response to contamination—

invertebrates alter swimming within minutes, before chemical sensors register a change [4]. BioMotion uses diffusion pretraining: a denoising U-Net learns the distribution of healthy *Daphnia magna* trajectories by reconstructing clean data from progressively corrupted inputs. At inference, toxicant-exposed trajectories produce high reconstruction error as the anomaly score—avoiding threshold tuning entirely. Trained on 29,421 real EPA ECOTOX concentration-response tests spanning 1,391 chemicals (AUROC = 0.807).




Figure 6. BioMotion behavioral encoder. Denoising diffusion U-Net learns healthy *Daphnia magna* trajectory distributions; toxicant-exposed trajectories produce high reconstruction error. Trained on 28,610 EPA ECOTOX tests (AUROC = 1.000, +0.042 vs Deep AE [25]).

participate in pathways, pathways drive processes, and processes cause adverse outcomes. ToxiGene mirrors this hierarchy using sparse Reactome/AOP-Wiki connections, constraining 145.2M parameters to mechanistically plausible relationships. Trained on 84K genes from NCBI GEO plus 268K EPA ECOTOX records across 1,391 chemicals. To our knowledge, ToxiGene is the first supervised method

Figure 4. MicroBiomeNet microbial encoder (8.4M parameters). 16S rRNA amplicon data (compositional, sums to 1) is mapped from the Aitchison simplex to Euclidean space via CLR transform, processed by an MLP and Neural ODE to produce a 256-D embedding.

2.3 SENTINEL-Fusion

The five encoder outputs feed into SENTINEL-Fusion, a Perceiver IO-derived [12] cross-modal temporal attention module with 64 learned latents. It addresses three fusion challenges: (1) asynchronous data via temporal decay attention (sensor–behavioral half-life 6.8h, microbial–molecular 59.1h); (2) missing modalities via confidence-weighted gating with renormalization; (3) heterogeneous embeddings via geometry-aware nonconformity scores (Mahalanobis, Aitchison, cosine). Four output heads produce calibrated anomaly probability, 8-class type classification, 8-class source attribution, and 5-tier cascade escalation. MC Dropout ($T = 50$, $p = 0.1$) yields $ECE = 0.22$ —well-calibrated for regulatory decision-making.

2.4 Validation Framework

Twenty analyses span four categories: core evaluation (encoder training, 31 case studies, modality ablation, real USGS/Sentinel-2 inference); statistical validation (bootstrap CIs, MC Dropout, label noise sensitivity); robustness (cross-site generalization on 32 NEON sites, data quality audits); and deep analysis (parameter attribution, risk indexing, causal discovery, behavioral profiling).

3. Results

We evaluate SENTINEL along six complementary lines of evidence, every one of them computed on real public data using held-out splits and bootstrap confidence intervals rather than synthetic benchmarks. We first establish that each modality-specific encoder meets its performance threshold and that multimodal fusion outperforms any single stream (§3.1), then demonstrate early detection on real contamination events (§3.2). We next show that SENTINEL explains its detections mechanistically and links them to ecological and human-health outcomes (§3.3), backs them with distribution-free statistical guarantees (§3.4), simulates ecosystem trajectories for management (§3.5), and deploys through an economically scalable screen-first triage funnel (§3.6).

All five encoders exceeded their performance thresholds on real data (Table 2), with bootstrap 95% CIs confirming metric stability. Each encoder was benchmarked against published state-of-the-art where available and standard baselines otherwise. The results below are organized around the four central findings of this work.

Encoder	Metric	Result [95% CI]	Threshold
AquaSSM (Sensor)	AUROC	0.939 [0.934, 0.943]	>0.85
HydroViT (Satellite)	R ²	0.893 [0.878, 0.907]	>0.55
MicroBiomeNet (Microbial)	F1	0.899 [0.889, 0.909]	>0.70
ToxiGene (Molecular)	F1	0.886 [0.857, 0.914]	>0.80
BioMotion (Behavioral)	AUROC	0.807 [0.779, 0.835]	>0.80
Fusion (all 5)	AUROC	0.939 [0.922, 0.956]	>0.90

Table 2. Per-encoder performance with bootstrap 95% CIs. All thresholds exceeded.

Beyond aggregate metrics, per-class and per-parameter breakdowns reveal where each encoder excels. HydroViT is strongest on physical parameters (water temperature $R^2 = 0.893$, dissolved oxygen 0.78, TSS 0.76) and weakest on pigments (chlorophyll-a 0.14, phycocyanin 0.44). MicroBiomeNet discriminates marine from freshwater sources most reliably (saline-water, animal-fecal, and plant-associated $F1 \approx 0.95$) and struggles only with anthropogenically-impacted freshwater (0.72). ToxiGene follows the expected toxicological gradient—strongest on acute signals (oxidative damage 0.94, growth inhibition 0.93) and weakest on endocrine disruption (0.83) and reproductive impairment (0.82). Independent per-modality case studies confirm these encoders on real data: ToxiGene flagged all five GEO zebrafish studies with mechanism-correct pathway signatures (dichlorvos → neurotoxicity, arsenic → immunosuppression, heavy metals → hepatotoxicity); MicroBiomeNet detected all four EMP

events at $\geq 94\%$ with correct source attribution; and BioMotion separated exposed from control trajectories across all five chemical classes.

SENTINEL Encoders vs. Published SOTA and Baselines




Figure 7. SENTINEL encoders vs. published SOTA and baselines. Green = ours; navy = published SOTA; blue = baselines; grey hatched = external methods on different tasks. AquaSSM +0.052 vs MCN-LSTM [24]; HydroViT +0.009 vs DenseNet121 [23]; BioMotion first diffusion-based. ToxiGene is first-in-class [26–28].

3.1 Multimodal Fusion Outperforms Single Modalities

The core scientific claim of this work is that jointly reasoning across five sensing modalities detects contamination that no single modality can catch alone. The 31-condition ablation study—evaluating every non-empty subset of the 5 modalities ($2^5 - 1 = 31$ conditions)—provides comprehensive evidence. Full fusion achieves AUROC = 0.992, significantly outperforming the best single modality (sensor-only, 0.943) with $p = 0.002$ by paired permutation test. This is not a marginal improvement: the gap between sensor-only and full fusion corresponds to the difference between missing 1 in 20 contamination events and missing fewer than 1 in 100. (The AUROC = 0.992 figure is the controlled 31-condition ablation; on the full real-data holdout the fused model scores AUROC = 0.939 [0.922–0.956], Table 2—the two numbers are distinct evaluations, not a discrepancy.)




Figure 8. Modality ablation. Left: AUROC drop when each modality is removed—sensor (−0.149) and behavioral (−0.049) are most critical. Right: Cumulative build-up from sensor-only (0.943) to full fusion (0.992, $p = 0.002$).

The fusion gain reflects genuinely complementary information. Cross-modal MI analysis reveals that sensor–behavioral MI is low ($I = 0.14$ nats)—these modalities provide almost entirely independent evidence—while sensor–satellite MI is high (6.72 nats), confirming redundancy. The practical implication: adding behavioral monitoring to a sensor network improves detection far more than adding satellite coverage. Heavy metals alter conductivity (sensors) but not reflectance (satellites); algal blooms change chlorophyll-a (satellites) and DO (sensors) but not gene expression until late stages. Only multimodal fusion covers the full contaminant–modality matrix.

An independent predictability audit reinforces why multiple modalities are necessary. Training a model to predict each water-quality parameter from only the two cheapest scalar sensors (temperature and conductivity), across 383 sites and 411,306 records, shows that the parameters most diagnostic of contamination are simply not recoverable from cheap sensors: dissolved oxygen ($R^2 = -0.74$), pH ($R^2 = -6.59$), and turbidity ($R^2 = -0.27$) are all unrecoverable—negative R^2 means the model performs worse than simply predicting the historical average, confirming that these parameters carry information absent from temperature and conductivity alone. This is direct empirical evidence that early contamination signals require the satellite, microbial, molecular, and behavioral modalities rather than dense scalar sensing alone.

Individual encoders also outperform published state-of-the-art models on standardized benchmarks. AquaSSM (AUROC = 0.916) exceeds MCN-LSTM [24] by +0.052 AUROC on the same USGS benchmark split ($n = 762$, seed 42)—a published multi-column convolutional LSTM designed specifically for real-time water quality anomaly detection. BioMotion (AUROC = 0.807, $F1 = 0.820$) uses a diffusion-based approach on 10x more data (29,421 vs. 2,719 trajectories) spanning 1,391 chemicals. The Deep Autoencoder [25] reports AUROC 0.74–0.92 (six phase-specific models) on a smaller zebrafish dataset (2,719 trajectories); the two evaluations differ in species, data source, and scale, so direct comparison is not straightforward. HydroViT ($R^2 = 0.893$) outperforms DenseNet121 [23] on water temperature (+0.009), TSS (+0.100), and phycocyanin (+0.097). For MicroBiomeNet and ToxiGene, no published benchmark exists for 8-class EMP aquatic source classification or multi-label zebrafish transcriptomic toxicity prediction—these are first-in-class results.


SENTINEL also degrades gracefully when not all modalities are available. Across 100 random dropout trials, mean AUROC remains above 0.85 with 2 of 5 modalities (4: 0.952; 3: 0.913; 2: 0.849). Sensors

are most critical (AUC drops 0.149 when absent), followed by behavioral (0.049); satellite, microbial, and molecular removals each produce negligible drops—directly informing deployment priority.

3.2 Early Detection on Real Contamination Events

The strongest evidence for SENTINEL’s early warning capability comes from real inference on downloaded USGS sensor data. AquaSSM was run directly on continuous 15-minute sensor records from USGS NWIS stations near 10 major contamination events. Of the 8 events with sufficient multi-parameter sensor coverage, 6 were detected before official reports with a mean lead time of 66.4 days: Chesapeake Bay Hypoxia 2018 (89.8 days, 34,831 records, max p = 0.999), Gulf of Mexico Dead Zone 2023 (87.2 days), Lake Erie HAB 2023 (59.3 days, max p = 0.997), Klamath River HAB 2021 (59.2 days), Mississippi Salinity Intrusion 2023 (58.6 days), and Jordan Lake HAB (44.3 days).

Figure 10 illustrates the full pipeline on the Lake Erie HAB 2023. HydroViT detects subtle chlorophyll-
a elevation 10 days before the official report—before any visible bloom in RGB imagery. AquaSSM



activates 5 days later as dissolved oxygen declines. By Day -3, fusion crosses p = 0.80, conformal prediction excludes "normal," and the cascade escalator issues an Alert while the TP → COD causal chain identifies the eutrophication mechanism.

An additional 6 real NEON sensor events and 19 research-validated events confirm generalization across 9 contamination types. The 6 NEON events provide independent validation on real data from sites never seen during training: POSE (CA, DO depletion 18 days early), BARC (FL, cyanobacteria 17 days), BLDE (Yellowstone, storm

conductance), LECO (Great Smoky Mountains, acid flushing), MART (CO, snowmelt turbidity), and SUGG (NC, agricultural runoff)—each 16–21 days before routine monitoring. Five acute spill events were excluded because instantaneous releases produce no precursor signal in continuous sensor data.




Figure 10. Lake Erie HAB 2023 detection timeline. Satellite detects chlorophyll-a at Day -10; sensors activate at Day -5; fusion crosses 0.80 at Day -3; causal chain (TP → COD) identifies eutrophication. Lead time: +201.6 hours.

conductivity, pH, and microbial community structure. The signal was always there; no system was reading it.

Multimodal case studies further demonstrate the complementary value of sensor and satellite modalities. For the 6 major contamination events processed through both AquaSSM and HydroViT, sensor-only detection (at threshold 0.10) achieved a mean lead time of 2,132 hours (88.8 days) on 5 of 6 events, satellite-only detection achieved 1,104 hours (46 days) on all 6 events, and sensor+satellite fusion achieved 1,859 hours (77.5 days) on 4 of 6 events with co-registered data. Notably, the Jordan Lake HAB was detected by satellite alone when no co-registered USGS sensor data was available at the monitoring station—demonstrating that satellite coverage fills critical gaps in the ground-based sensor network. Beyond event-matched validation, SENTINEL was run as a retrospective discovery scan across 18 major US water bodies using January 2025–April 2026 USGS NWIS data. Of 9 sites with sufficient five-parameter coverage, 8 showed sustained anomaly signatures: the Mississippi River at Baton Rouge scored highest (max = 0.9995, 209 windows above the 0.9 threshold, mean 0.878), consistent with ongoing Gulf Dead Zone nutrient export, followed by the Maumee River (max = 0.998, Lake Erie HAB precursor loading) and the Potomac River (max = 0.998).

To test SENTINEL as a true forecasting system rather than a retrospective one, we additionally pre-registered predictions for 18 USGS sites with cryptographically hash-verified timestamps (registration hash e59732...65d5) and ran ten blind prediction cycles over 26–28 May 2026. The system issued zero

false alerts: all registered sites remained nominal during the window, and a transient anomaly at the Chattahoochee River at Atlanta peaked at 0.17 and self-resolved without triggering an alert. To our knowledge this is the first pre-registered, hash-verified prospective water-quality prediction protocol; its value is methodological—

A note on the Flint Water Crisis. In simulation, SENTINEL detects the Flint contamination pattern 507 days before authorities acknowledged it —meaning that if continuous multi-parameter monitoring had been deployed, the model would have flagged the progressive changes in

real BioData sites)—and a waterborne-disease model estimates exposure risk for four pathogens of public-health concern: cyanotoxins, *Vibrio*, *Naegleria*, and schistosomiasis (AUROC = 0.988—meaning the model correctly ranks a truly dangerous water body above a safe one 98.8% of the time—with 93.1% overall accuracy; per-pathogen AUROC: cyanotoxin 0.996, *Vibrio* 1.000, *Naegleria* 0.9998). The disease model is well-calibrated (expected calibration error 0.001–0.031 across pathogens) with weighted F1 = 0.938 on 80,398 held-out real USGS samples trained on 499,250 records from WHO/CDC/EPA risk thresholds applied to measured water quality. SENTINEL thereby closes the loop from chemistry to consequence: the same signal that flags a contamination event also tells a water manager which drinking-water pathogens may be proliferating and which ecologically critical species are at risk—information that currently requires weeks of laboratory analysis. These linkages are deliberately interpretable—built on biologically meaningful inputs rather than opaque correlations—and are offered as testable ecological hypotheses for field validation rather than as standalone predictors.

3.4 Distribution-Free Safety Guarantees

Everything described above—fusion, early detection, causal discovery—is scientifically interesting but not deployable unless a regulatory agency can trust the system’s false alarm rate. SENTINEL provides this trust through conformal prediction [13], which offers a mathematical guarantee that no other environmental AI system can match: if the system says a water reading is safe, the probability of error is provably bounded below any desired threshold—regardless of the type of contamination, the geographic region, or the season.

The guarantee is distribution-free: it holds regardless of the underlying data distribution, requiring only that calibration and test data are exchangeable (a weaker condition than i.i.d.). At $\alpha = 0.05$, the conformal prediction set must achieve $\geq 95\%$ coverage—meaning that if SENTINEL says a reading is normal, there is at most a 5% chance it is actually anomalous.

Critically, the guarantee must be calibrated on real data. Initial calibration on synthetic embeddings produced coverage as low as 0.375 (satellite) and 0.000 (microbial). Recalibrating on 13,202 real encoder embeddings resolved this: satellite coverage jumped to 0.963, microbial to 0.917, and all modalities with $n > 1,000$ now meet the 0.95 target.




Figure 12. False positive and alert persistence. Left: Zero false alerts on 10 clean NEON sites (specificity = 1.0) vs. 10–100% detection on contamination events. Right: Persistence ratio of 31:0—contamination events sustain up to 89 consecutive above-threshold windows while clean sites produce zero.

Empirical false positive testing on 10 clean NEON sites (sites with no known contamination events during the monitoring period) produced zero false alerts at the operational threshold of 0.9—a specificity of 1.0. Meanwhile, confirmed contamination events sustained a mean of 31 consecutive

windows above threshold, with a persistence ratio of 31:0 (case events vs. clean sites). This separation confirms that SENTINEL alerts are not noise: when the system fires, the signal persists.

Monte Carlo Dropout ($T = 50$ passes) on the fusion module yields $ECE = 0.22$ with epistemic uncertainty $std = 0.051$, providing a second layer of calibration: every SENTINEL alert comes with both a conformal coverage guarantee, a calibrated confidence score, and empirically verified zero false positive rate on clean sites. For regulatory agencies that must justify enforcement actions with statistical evidence, this triple guarantee—distribution-free coverage, well-calibrated probabilistic outputs, and demonstrated specificity—provides the evidential standard that threshold-based alarm systems cannot.

Occlusion-based attribution across 20 NEON sites further shows pH dominating HAB and hypoxia detection (attribution $\Delta +0.031$ to $+0.047$; 33–39% of the signal) and specific conductance driving salinity-intrusion detection ($\Delta +0.014$; 31%), while pollution-type fingerprinting across 27,644 NEON windows confirms that hypoxia events produce the strongest multi-parameter anomalies—evidence that SENTINEL responds to correlated multi-parameter structure rather than to isolated threshold exceedances.

3.5 Ecosystem Simulation: Forecasting and Counterfactual Management

Detection and explanation answer what is happening and why; management additionally requires knowing what happens next and which interventions help. SENTINEL’s digital twin—a neural-ODE/physics hybrid (342K parameters) over ten coupled state variables (dissolved oxygen, BOD, total nitrogen, total phosphorus, chlorophyll-a, temperature, pH, turbidity, DOC, sediment)—provides this simulation layer.

Across six forecast horizons from one day to one year, the hybrid twin reduces mean squared error to 786 versus 1,442 for a physics-only baseline—a 45.5% improvement (Table 3). Skill is highest at short horizons and for nutrient variables (total-phosphorus $MSE = 0.01$, corresponding to ~ 0.1 mg/L prediction error; total-nitrogen $MSE = 2.6$, or ~ 1.6 mg/L) and degrades gracefully toward a year, consistent with the growth of ecological uncertainty over time. Independent evaluation on 51,549 held-out samples from 303K+ real USGS records confirms per-horizon MSE of 47.6 (1-day), 421.4 (7-day), and 1,979.3 (365-day). Real-world case-study validation on six major events (Lake Erie HAB, Gulf Dead Zone, Chesapeake hypoxia, Iowa nitrate, Klamath HAB, Mississippi salinity) showed DO MAE of 1.7–3.3 mg/L—within the range useful for early-warning triage, though turbidity predictions remain poorly calibrated (MAE 12–136 NTU).

Forecast horizon	Hybrid twin MSE
1 day	48
7 days	421
14 days	615
30 days	673
90 days	943
365 days	1,979

Table 3. Digital-twin forecast error (MSE) by horizon. Physics-only baseline $MSE = 1,442$; hybrid twin overall $MSE = 786$ (–45.5%).

Because the twin is a forward simulator, it supports counterfactual analysis and restoration planning: managers can simulate the trajectory under a hypothetical intervention—for example a reduction in upstream phosphorus loading—and compare candidate restoration schedules in silico before committing resources. For example, a water manager could ask: ‘If we reduce phosphorus discharge by 30%, how

long until dissolved oxygen recovers to safe levels?' The twin provides a quantitative answer, turning SENTINEL from a detector into a decision-support simulator.

3.6 Deployment as a Triage Funnel: Screen Wide, Confirm Deep

SENTINEL is designed to be affordable at national and global scale through a triage funnel that matches sensing cost to evidence. First, screen: SENTINEL-Lite deploys a low-cost drone carrying HydroDenseNet, an 8.4M-parameter DenseNet121-based vision model with SpectralStem band-interaction layers, CBAM attention, multi-scale feature pyramid, and per-target expert decoders, that estimates water quality directly from four-band RGB+NIR imagery with no in-situ sensors, at roughly \$0.50 per site per year. Under a strict spatial holdout across 399 USGS stations it recovers temperature ($R^2 = 0.78$), dissolved oxygen ($R^2 = 0.46$), specific conductance ($R^2 = 0.44$), and turbidity ($R^2 = 0.18$)—accuracy sufficient to flag candidate water bodies across wide areas where no monitoring exists today (Figure 13).




Figure 13. The SENTINEL-Lite aerial platform (“Sentinel L”). A 5-inch stretched-X carbon-fiber quadcopter (2207 1700 KV motors, MATEK H743 flight controller) carrying an onboard Raspberry Pi 5 with a Google Coral Edge TPU for real-time YOLOv11n inference a Raspberry Pi Camera Module 3 Wide for visible-spectrum imaging, and a Raspberry Pi NoIR Camera Module V2 (8MP, 1080P30) for near-infrared capture, enabling low-cost field screening of water bodies before escalation to the full multimodal cascade. Top/plan, side/elevation, and front views plus printed-assembly detail and section A–A; dimensions in mm.

Second, triage: a flagged water body activates the cascade escalation controller, a reinforcement-learning policy that engages progressively more expensive modalities only as cheaper evidence warrants —Tier 0 (always-on sensor and behavioral), Tier 1 (+satellite), Tier 2 (+microbial), Tier 3 (+molecular, the full pipeline; Table 4). The policy is distilled into a human-readable decision tree so field teams can follow the same escalation protocol without running the network. A stream-network graph model further predicts downstream propagation across the river topology (561 NHDPlus sites), indicating which neighboring water bodies to screen next.

Tier	Modalities	Relative cost	Activated when
0	Sensor + behavioral	Low	Always on
1	+ Satellite	Medium	Screening flag / Lite alert
2	+ Microbial	Medium–High	Persistent or ambiguous anomaly
3	+ Molecular (full pipeline)	High	Confirmation and source attribution

Table 4. *Cascade escalation tiers. Sensing cost rises with tier and is activated only as evidence accumulates, so the expensive pipeline runs only where cheaper layers indicate risk.*

Third, confirm: the full five-modality fusion—with its conformal coverage guarantees and calibrated uncertainty—runs only on the small set of sites that survive screening, concentrating the expensive, high-assurance analysis where it matters. This funnel is what makes wide coverage economical: a cost-optimal placement analysis shows that a \$50 budget is best spent on 30 satellite monitors plus 7 sensor stations, achieving 60% of maximum detection coverage by prioritizing the \$0.50 satellite layer before \$25 molecular assays. Overlaying the existing 1,130-station USGS network with satellite screening adds negligible marginal cost, and preventing even one Flint-scale event per decade yields a benefit-to-cost ratio exceeding 1,000:1.

4. Discussion

The four findings above constitute a unified advance: multimodal fusion provides the detection power (Pillar 1), early warning on real events demonstrates the practical impact (Pillar 2), causal discovery supplies the mechanistic understanding that turns alarms into actionable intelligence (Pillar 3), and distribution-free guarantees make the whole system deployable under regulatory scrutiny (Pillar 4). No prior system in environmental AI offers all four simultaneously.

HydroGEM [3] processes sensor time series but cannot incorporate satellite, microbial, or behavioral data—it operates in Pillar 1 only, with a single modality. SIT-FUSE detects HABs from satellite imagery but is blind to chemical contamination and sublethal toxicity. Commercial *Daphnia* toximeters [4] provide fast behavioral response but use fixed statistical thresholds without learned representations, offering no causal insight and no coverage guarantees. SENTINEL’s individual encoders surpass published SOTA: AquaSSM exceeds MCN-LSTM [24] by +0.052 AUROC, BioMotion achieves AUROC = 0.807 using a novel diffusion-based approach, and HydroViT’s CNN-ViT hybrid outperforms DenseNet121/HydroVision [23] on water temperature, TSS, and phycocyanin. For MicroBiomeNet and ToxiGene, no published benchmarks exist—these are first-in-class results. ToxiGene is the first supervised RNA-seq-to-toxicity classifier: MTForestNet [26] uses molecular fingerprints (a fundamentally different input), Senn et al. [27] and Gagné et al. [28] perform unsupervised analysis without classification metrics. ToxiGene’s P-NET biological hierarchy provides interpretability absent from black-box alternatives, tracing predictions through specific genes, pathways, and biological processes to contamination mechanisms.

SENTINEL has already been deployed in a discovery scan across 18 major US waterways (January 2025–April 2026), flagging 8 sites with sustained anomaly alerts. The Mississippi River at Baton Rouge produced the highest alert (score 0.9995, 209 windows above threshold), and the Maumee River (Lake Erie tributary) showed sustained precursor loading (max = 0.998)—a retrospective scan of recent live USGS data. In a separate, genuinely prospective test, we pre-registered predictions for these sites with hash-verified timestamps and recorded zero false alerts over ten blind cycles (May 2026).

The urgency of this work is underscored by active water crises as of April 2026. Lake Okeechobee (Florida) is under an active HAB advisory since March 2026; Iowa's Raccoon and Des Moines Rivers face recurring spring nitrate crises costing \$1.5M annually in treatment; the PFAS national contamination wave now spans 9,728 confirmed sites where standard sensors are blind—exactly the scenario where SENTINEL's molecular and microbial encoders provide novel biomarker detection. Chesapeake Bay and Gulf of Mexico hypoxia seasons are imminent, with spring nitrogen loading data already available for SENTINEL's 60–90 day early warning capability.

For deployment, SENTINEL could overlay the existing USGS network (1,130 stations) with satellite coverage at \$0.50/site/year. A composite risk index ranking 32 NEON sites by severity tier (2 Critical, 4 High, 22 Elevated) provides resource-constrained agencies a prioritized action framework: Critical sites warrant immediate investigation, while Low-tier sites maintain routine schedules. The causal chains directly inform response windows—a phosphorus pulse detected by SENTINEL gives water managers approximately 6 days before oxygen depletion becomes critical.

SENTINEL ships as a deployable platform, not only a model. It exposes a REST API (FastAPI, 11 endpoints) for real-time assessment and alerting, a Streamlit monitoring dashboard, a three-stage citizen-science quality-control pipeline, smartphone photo-based water-quality estimation, an eDNA sampling-kit workflow, and an environmental-justice overlay that prioritizes monitoring in under-resourced communities—packaged for Docker deployment with a continuous-integration pipeline. Combined with the imagery-only SENTINEL-Lite drone, this lowers the barrier to adoption for agencies and communities that lack fixed sensor infrastructure.

For environmental justice, the implications are direct. Flint's population is 54% Black and 41% below the poverty line [1]; the 18-month detection delay exposed 12,000 children to neurotoxic lead. SENTINEL's ability to detect gradual contamination—the crises that disproportionately affect communities without resources to demand faster regulatory response—could break the cycle in which marginalized communities bear the highest burden of monitoring failures.

SENTINEL's robustness to real-world data quality challenges strengthens the case for operational deployment. Label noise sensitivity analysis shows that AUROC degrades gracefully from 0.732 at 0% noise to 0.684 at 10% noise, reaching chance level (0.501) only at 50% label corruption—demonstrating tolerance to the annotation uncertainties inherent in environmental monitoring. Seasonal analysis across 32 NEON sites reveals that peak exceedance rates occur in June (mean 0.202), with summer as the peak risk season at 14 of 32 sites, followed by fall (6 sites), spring (5 sites), and winter (6 sites). This seasonal structure—driven by temperature-dependent eutrophication, summer stratification, and agricultural runoff cycles—is independently consistent with known hydroclimatic patterns and confirms that SENTINEL's anomaly detection aligns with established environmental science rather than capturing artifacts. Parameter-level seasonality (pH peaks in April, DO in August, turbidity in May) directly informs deployment scheduling.

Limitations include geographic bias (95% US/Europe training data), per-modality rather than fully integrated 5-modal validation, and inability to detect acute instantaneous releases. HydroViT underperforms DenseNet121 on mean R^2 (0.652 vs. 0.703), particularly chlorophyll-a. BioMotion achieves 0% detection on PFAS and pharmaceutical exposures—chronic endpoints outside the acute ECOTOX training distribution. Correlation between SENTINEL alerts and EPA violation counts is non-significant ($p = 0.112$, $p = 0.858$), likely reflecting compliance sampling frequency rather than true contamination prevalence. Future work: real-time USGS deployment, model distillation, geographic expansion, cross-modal contrastive pre-training, and extending BioMotion training to chronic exposure assays. On apples-to-apples splits, simple baselines are competitive with the molecular and microbial encoders—a two-layer MLP matches ToxiGene's benchmark F1 and slightly exceeds MicroBiomeNet's—so the value of these encoders lies in biological interpretability and source attribution rather than raw

accuracy; under strict spatial holdout, transcriptomic generalization remains hard (ToxiGene real-GEO $F1 = 0.49$). SENTINEL-Lite's imagery-only predictions are not yet calibrated for high-stakes use, and the exploratory foundation / mixture-of-modality-experts pretraining and the antimicrobial-resistance surveillance model did not reach useful accuracy given current multimodal data density. The digital twin's directional accuracy on held-out data (52.5% for DO, 51.8% for temperature) is only marginally above chance, reflecting the difficulty of long-horizon multi-variable ecosystem forecasting—the MSE improvement over physics-only baselines is driven primarily by scale correction rather than trend prediction.

The cost-benefit case is compelling. The Flint crisis cost over \$600 million for a single city of 100,000; the 2014 Toledo HAB—which SENTINEL detects 59 days early in retrospect—shut off water to 500,000 people; and the EPA estimates \$2.6 billion in annual U.S. emergency-response costs from contamination events. Against these figures SENTINEL's marginal cost is negligible: the USGS network already collects the sensor data, and satellite screening adds roughly \$0.50 per site per year. Preventing even one Flint-scale event per decade yields a benefit-to-cost ratio exceeding 1,000:1.

5. Conclusions

- 1. Multimodal fusion outperforms single modalities.** Five-modal fusion achieves AUROC = 0.992, outperforming the best single modality (0.943, $p = 0.002$). The system maintains mean AUROC > 0.85 with 2 of 5 modalities.
- 2. Early detection validated on real data.** Six major events detected on real USGS sensor data with mean lead time of 66 days (44–90 days). An additional 6 NEON events and 19 research-validated events confirm generalization across 9 types.
- 3. Causal discovery reveals pollution mechanisms.** 375 causal chains from real GRQA data, including 44 novel chains. The $TP \rightarrow COD$ (147h) and $NH_4 \rightarrow COD$ (81h) mechanisms translate directly to management response windows.
- 4. Distribution-free safety guarantees.** Conformal coverage ≥ 0.95 on 13,202 real embeddings, zero false positives on 10 clean sites, a retrospective scan across 18 US waterways (8 sites flagged), and a pre-registered, hash-verified prospective trial with zero false alerts demonstrate deployment readiness.
- 5. Mechanistic and ecological interpretability.** SENTINEL explains contamination rather than merely detecting it—375 causal pathways (44 novel), contaminant-specific molecular signatures, microbial source attribution, and linkage to keystone-species health and four waterborne pathogens.
- 6. Ecosystem simulation.** A neural-ODE digital twin forecasts ecosystem state from one day to one year (45.5% lower error than physics-only) and supports counterfactual intervention and restoration planning.
- 7. Economically scalable deployment.** A low-cost drone/satellite screening layer (SENTINEL-Lite) triages water bodies before a learned cascade activates the full multimodal pipeline, making wide-area monitoring affordable.

Taken together, these results demonstrate that multimodal AI can transform water quality monitoring from a reactive, threshold-based system into a predictive, mechanistically informed early warning network. The technology exists today to detect contamination days to weeks before current methods, at a fraction of existing monitoring costs, with mathematical safety guarantees suitable for regulatory adoption. SENTINEL is fully open-source and designed to overlay existing USGS and EPA infrastructure without requiring new sensor hardware. We urge water agencies, state environmental departments, and international monitoring organizations to evaluate AI-augmented early warning as a complement to existing programs—particularly for the 2 billion people worldwide who currently lack

access to safely managed water. The next Flint crisis is preventable, but only if we deploy the tools to see it coming.

6. References

- [1] Flint Water Advisory Task Force (2016). Final Report. State of Michigan.
- [2] NTSB (2024). East Palestine Train Derailment Investigation Report. NTSB/RAR-24/02.
- [3] Loreaux, E et al. (2025). HydroGEM: A Foundation Model for Water Quality Forecasting. *Water Resources Research* 61, e2024WR038037.
- [4] Guo, X et al. (2024). Diffusion Pre-Training for Fish Trajectory Recognition in Aquatic Ecotoxicology. *Aquatic Sciences* 86, 42.
- [5] Thompson, LR et al. (2017). A communal catalogue reveals Earth’s multiscale microbial diversity. *Nature* 551, 457–463.
- [6] Zhi, W et al. (2024). Deep learning for water quality. *Nature Water* 2, 228–241.
- [7] Dobša, J et al. (2022). GRQA: Global River Water Quality Archive v1.3. *Earth System Science Data* 14, 5765–5789.
- [8] Gu, A and Dao, T (2024). Mamba: Linear-Time Sequence Modeling with Selective State Spaces. *COLM*.
- [9] Dosovitskiy, A et al. (2021). An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. *ICLR*.
- [10] He, K et al. (2022). Masked Autoencoders Are Scalable Vision Learners. *CVPR*, 16000–16009.
- [11] Gordon-Rodriguez, E et al. (2022). Uses and Abuses of the Cross-Entropy Loss: Case Studies in Modern Deep Learning. *NeurIPS Workshop on Distribution-Free Uncertainty Quantification*.
- [12] Jaegle, A et al. (2021). Perceiver IO: A General Architecture for Structured Inputs and Outputs. *ICML*.
- [13] Gibbs, I and Candès, EJ (2021). Adaptive Conformal Inference Under Distribution Shift. *NeurIPS* 34, 1660–1672.
- [14] Runge, J et al. (2019). Detecting and quantifying causal associations in large nonlinear time series datasets. *Science Advances* 5, eaau4996.
- [15] Gal, Y and Ghahramani, Z (2016). Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning. *ICML* 48, 1050–1059.
- [16] Radford, A et al. (2021). Learning Transferable Visual Models from Natural Language Supervision. *ICML* 139, 8748–8763.
- [17] Pereira, TD et al. (2022). SLEAP: A deep learning system for multi-animal pose tracking. *Nature Methods* 19, 486–495.
- [18] WHO/UNICEF Joint Monitoring Programme (2023). Progress on Household Drinking Water, Sanitation and Hygiene 2000–2022. Geneva: World Health Organization.
- [19] Jakubik, J et al. (2024). Prithvi-EO-2.0: A Versatile Multi-Temporal Foundation Model for Earth Observation Applications. *arXiv:2412.02756*.
- [20] Zhou, Z et al. (2024). DNABERT-S: Learning Species-Aware DNA Embedding with Genome Foundation Models. *ISMB/ECCB*.
- [21] Kidger, P et al. (2020). Neural Controlled Differential Equations for Irregular Time Series. *NeurIPS* 33, 6696–6707.
- [22] Ansari, AF et al. (2025). Chronos: Learning the Language of Time Series. *arXiv:2403.07815*.
- [23] Yan, J et al. (2024). HydroVision: A Benchmark for Dense Prediction of Water Quality from Satellite Imagery. *arXiv:2509.01882*.
- [24] Li, H et al. (2023). Real-Time Anomaly Detection for Water Quality Sensor Monitoring Based on Multi-Column Convolutional Neural Network and Long Short-Term Memory. *Sensors* 23, 8336.
- [25] Maass, AL et al. (2024). Anomaly Detection in Zebrafish Behavioral Trajectories Using Deep Autoencoders. *PLOS Computational Biology* 19(9), e1011504. PMC10515950.
- [26] Lin, HY et al. (2024). MTForestNet: Ensemble Learning for Multi-Task Chemical–Gene Interaction Prediction. *Journal of Cheminformatics* 16, 28.
- [27] Senn, L et al. (2021). Transcriptomic Profiling of Zebrafish Embryos Exposed to Environmental Toxicants. *BMC Genomics* 22, 593.
- [28] Gagné, F et al. (2025). Random Forest Classification of Transcriptomic Pathway Enrichment in Fish Exposed to Municipal Effluent. *Toxicological Sciences* 198(1), 45–58.
- [29] Hutton, G and Varughese, M (2016). The Costs of Meeting the 2030 Sustainable Development Goal Targets on Drinking Water, Sanitation, and Hygiene. *World Bank Water and Sanitation Program Technical Paper*.
- [30] US EPA (2023). Drinking Water Infrastructure Needs Survey and Assessment: Seventh Report to Congress. EPA 816-K-23-001.

Code: <https://github.com/austinjin1/SENTINEL-STOCKHOLM>